



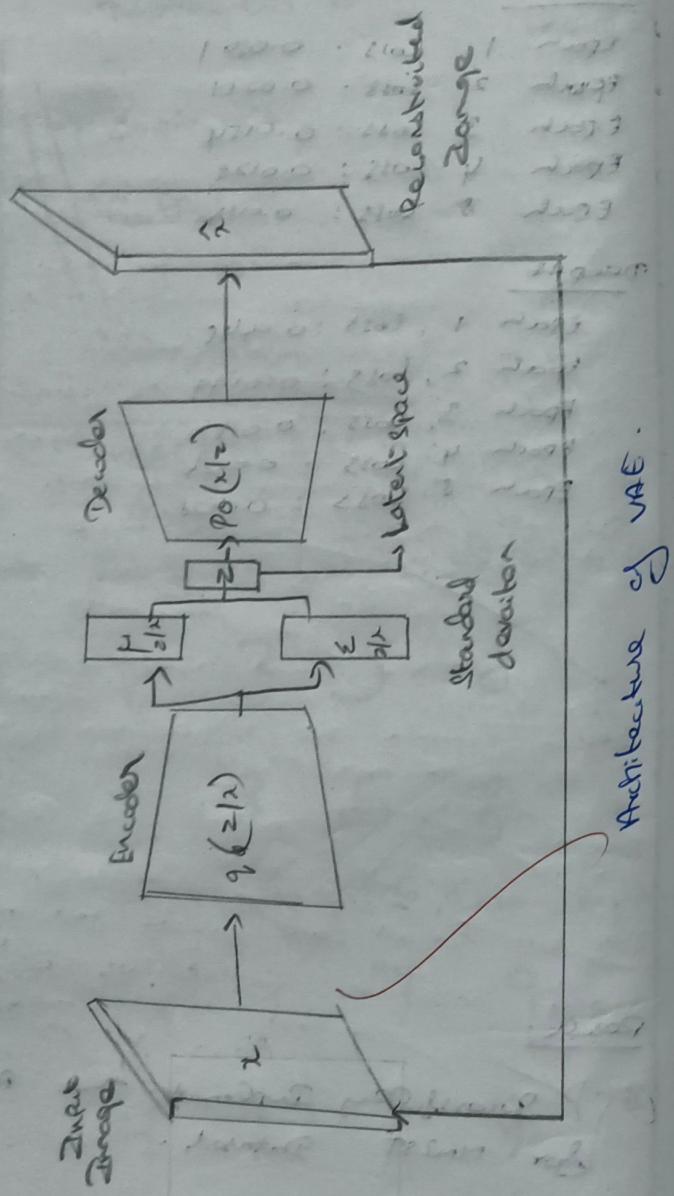
Superior®
Note Books

Index

Name: ... Vishwanath S
Subject: ... DL
Std.: Sec.: Roll No.
School:

No.	Date	Title	Marks	Signature
1	24/7/25	Exploring the DL platforms.		effortless
2	31/7/25	Implement a classifier using open source dataset		
3	7/8/25	Study the classifier with respect to statistical parameters.		effortless
4	14/8/25	Build a simple feed forward neural network to recognize handwritten characters.		effortless (12/8/25)
5	22/8/25	Study of Activation functions & their role.		effortless (27/9/25)
6	9/9/25	Implement Convolutional Descent & Backpropagation in Deep neural Network		effortless (21/9/25)
7	16/9/25	Build a CNN model to classify cat > dog image.		effortless (12/9/25)
8	23/9/25	Build a RNN network.		effortless (30/9/25)
9	9/10/25	Experiment using LSTM		effortless
10	9/10/25	Autoencoder		effortless
11	9/10/25	VARIATIONAL AUTOENCODER		effortless
12	17/10/25	Deep convolutional GAN		effortless

9/10/2023
Lab 11: Experiments using Variational Autoencoder (VAE)



Aim:

To implement a Variational Autoencoder to learn probabilistic latent representations of MNIST image and generate new images from the learned latent space.

Objectives

- 1) To understand the difference between Autoencoders & variational Autoencoder.
- 2) To apply probabilistic encoding using mean & variance.
- 3) To generate new images by sampling from the latent space.
- 4) To visualize the learned latent distribution & reconstructed images

Observation

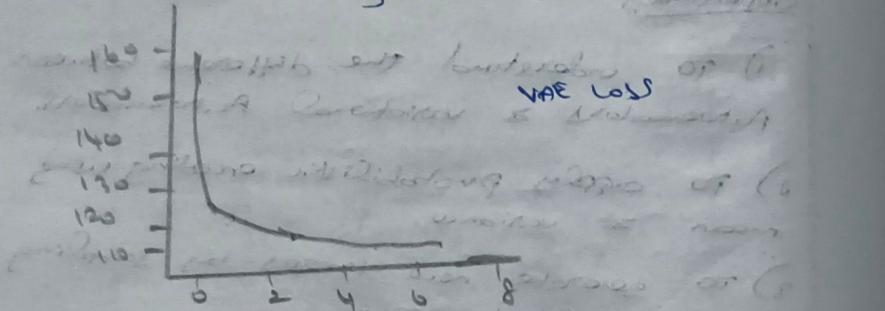
- The VAE extends the Autoencoder by introducing probabilistic latent variables
- It learns a distribution rather than a fixed encoding, allowing sampling & generation of new data.
- The model is trained using a reconstruction loss and a KL divergence loss to balance accuracy & latent regularization
- The generated samples resemble hand written digits, showing that the model learned the data distribution.

Output:

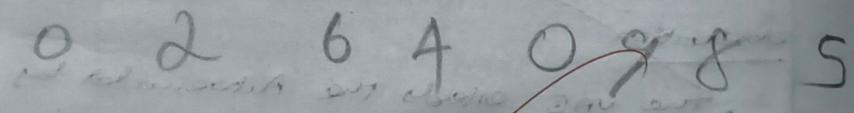
Epoch 1 loss : 164.378
 Epoch 2 loss : 121.282
 Epoch 3 loss : 115.114, 521
 Epoch 4 loss : 111.603
 Epoch 5 loss : 109.8342

Epoch 6 loss : 108.6767
 Epoch 7 loss : 107.829
 Epoch 8 loss : 107.117
 Epoch 9 loss : 106.6916
 Epoch 10 loss : 106.2222

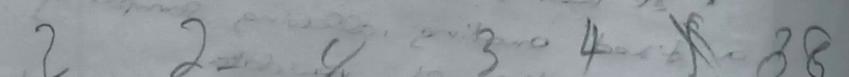
Training loss of VAE



TOP : original images | Bottom : VAE Reconstruction



Generated Digits from Random samples from Latent Space



and also some images are generated which are not present in the original dataset.



These images show that the VAE has learned to generate new images that are similar to the training data.

Pseudo code:

BEGIN

Step 1: Import Libraries

Step 2: Load & normalize MNIST dataset

Step 3: Define encoder

- Input: 784-dimensional image

- Dense hidden layers

- Output: mean (μ) and log variance (σ^2)

Step 4: Sample latent vector $z = \mu + \sigma \epsilon$

Step 5: Define decoder

- Input: z

- Dense layers to reconstruct 784-dimensional output

Step 6: Define loss function

- Reconstruction loss + KL divergence

Step 7: Train VAE using training data

Step 8: Generate new images by sampling from latent space

Step 9: Visualize reconstructed and generated images

END.

Result:

Successfully implemented the VAE for MNIST dataset.

CO dltlab11.ipynb ⭐ 🌐

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all Connect T4

pip install torch torchvision matplotlib

```
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages (2.8.0+cu126)
Requirement already satisfied: torchvision in /usr/local/lib/python3.12/dist-packages (0.23.0+cu126)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch) (3.20.0)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12/dist-packages (from torch) (4.15.0)
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch) (55.2.0)
Requirement already satisfied: sympy>1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch) (1.13.3)
Requirement already satisfied: networkx in /usr/local/lib/python3.12/dist-packages (from torch) (3.5)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from torch) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.12/dist-packages (from torch) (2025.3.0)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.80)
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages (from torch) (9.10.2.21)
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.4.1)
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages (from torch) (11.3.0.4)
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages (from torch) (10.3.7.77)
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages (from torch) (11.7.1.2)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages (from torch) (12.5.4.2)
Requirement already satisfied: nvidia-cusparselt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages (from torch) (0.7.1)
Requirement already satisfied: nvidia-nccl-cu12==2.27.3 in /usr/local/lib/python3.12/dist-packages (from torch) (2.27.3)
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages (from torch) (12.6.85)
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages (from torch) (1.11.1.6)
Requirement already satisfied: triton==3.4.0 in /usr/local/lib/python3.12/dist-packages (from torch) (3.4.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from torchvision) (2.0.2)
Requirement already satisfied: pillow!=8.3.*,>5.3.0 in /usr/local/lib/python3.12/dist-packages (from torchvision) (11.3.0)
Requirement already satisfied: contourpy==1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.9)
Requirement already satisfied: packaging==20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (25.0)
Requirement already satisfied: pyparsing==2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.2.5)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (2.9.0.post0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7->matplotlib) (1.17.0)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages (from sympy>1.13.3->torch) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages (from jinja2>=torch) (3.0.3)
```

vae_mnist.py
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader

CO ditlab11.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▾

Connect 14 ▾ V

```
11
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt

# ---- Device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# ---- Load MNIST dataset
tf = transforms.Compose([transforms.ToTensor()])
train_data = datasets.MNIST(root='data', train=True, transform=tf, download=True)
train_loader = DataLoader(train_data, batch_size=128, shuffle=True)

# ---- Variational Autoencoder
class VAE(nn.Module):
    def __init__(self):
        super(VAE, self).__init__()
        self.fc1 = nn.Linear(784, 400)
        self.fc_mu = nn.Linear(400, 20)      # Mean of latent vector
        self.fc_logvar = nn.Linear(400, 20) # Log variance of latent vector
        self.fc2 = nn.Linear(20, 400)
        self.fc3 = nn.Linear(400, 784)
        self.relu = nn.ReLU()
        self.sigmoid = nn.Sigmoid()

    def encode(self, x):
        h1 = self.relu(self.fc1(x))
        mu = self.fc_mu(h1)
        logvar = self.fc_logvar(h1)
        return mu, logvar

    def reparameterize(self, mu, logvar):
        std = torch.exp(0.5 * logvar)
        eps = torch.randn_like(std)
        return mu + eps * std # z = μ + σ * ε

    def decode(self, z):
        h2 = self.relu(self.fc2(z))
        return self.sigmoid(self.fc3(h2))

    def forward(self, x):
        x = x.view(-1, 784)
        mu, logvar = self.encode(x)
```

CO dltlab11.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all ▶

Connect T4 ▶

```
# ---- Loss Function (Reconstruction + KL Divergence)
def vae_loss(recon_x, x, mu, logvar):
    BCE = nn.functional.binary_cross_entropy(recon_x, x.view(-1, 784), reduction='sum')
    # KL Divergence: how latent distribution differs from standard normal
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    return BCE + KLD

# ---- Model, Optimizer
model = VAE().to(device)
optimizer = optim.Adam(model.parameters(), lr=1e-3)

# ---- Training
epochs = 10
train_losses = []

for epoch in range(1, epochs + 1):
    model.train()
    total_loss = 0
    for imgs, _ in train_loader:
        imgs = imgs.to(device)
        optimizer.zero_grad()
        recon, mu, logvar = model(imgs)
        loss = vae_loss(recon, imgs, mu, logvar)
        loss.backward()
        total_loss += loss.item()
        optimizer.step()

    avg_loss = total_loss / len(train_loader.dataset)
    train_losses.append(avg_loss)
    print(f"Epoch [{epoch}/{epochs}] Loss: {avg_loss:.4f}")

# ---- Plot Training Loss
plt.plot(train_losses, label="VAE Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training Loss of Variational Autoencoder")
plt.legend()
plt.show()

# ---- Reconstruction Visualization
model.eval()
imgs, _ = next(iter(train_loader))
imgs = imgs[:8].to(device)
with torch.no_grad():
    recons, _, _ = model(imgs)
```

Variables Terminal

